

Q. Suppose you ran import in db & import is running slow, how you can speed up process.

① Parallelism :- data pump support parallel processing.

impdp dumpfile=d1.dmp parallel=4.

② Buffer size :- increase the buffer size.

impdp dumpfile=d1.dmp data_options=Buffer Size: 1024

③ Disable index & constraints :-

impdp dumpfile=d1.dmp transform=disable_constraints,
--> indexes;

④ use direct path:-

impdp dumpfile=d1.dmp Direct=Y ...

Q. Can I increase filesize while running parallel.

Yes.

When you running data pump operation in 11e you may encounter this scenario.

To increase the file size of each 11e process, you can adjust the filesize parameter.

exp impdp dumpfile=d1.dmp parallel=4 Filesize=500m.

Q. When facing a critical activity at the end of my shift, what will your approach. handover or continue. → depend on nature of activity, its urgency & the potential impact on the project. Then I working continue

Q. What do you think, db activity is small?

No.

db activity is very responsible work. because we handle or check ① transaction volume ② data size

③ complexity of queries. ④ system performance

⑤ Data changes ⑥ application workload

⑦ Business impact.

Q. Tell me any occasion where you taken ownership of completed activity on your own?

- Q. Explain Startup Sequence of a Cluster?
- Startup
- ① stop db ←
 - ② stop ggs
 - ③ crsctl stop crs
 - ④ if oracle listener is up then stop manually.
 - ⑤ check status.
- Shut down
- ① cluster are normally up
 - ② if not up then use crsctl start crs
 - ③ Start all listeners
 - ④ start ggs
 - ⑤ start db.

Q. do you know export & import? → utility
yes.

Export: is used to extract data & db objects from an oracle db & store them in a binary.

Import: is used to read data & db object from an export dump file created by export utility & load them into an oracle db.

- Q. if you are refreshing S Schema at a time & source name & target name tablespace Schema are different, what you do?
- ① create mapping betw Source & target tablespace.
 - ② export data with tablespace mapping.
`expdp dumpfile=d1.dmp remap_tablespace=source_t1:target_t1, source_t2:target_t2.`
 - ③ transfer & import data:
`impdp dumpfile=d1.dmp remap_tablespace=source_t1:target_t1, source_t2:target_t2;`

- Q. why remap need to be used?
it allows you to modify object names, tablespace & other attributes during export & import.
- use of remap:- different tablespace during import
- ① Remap_tablespace :- allow you to remap object to a different tablespace.
 - ② Remap_schema :-

encrypting Rman backups:

- i) Connect with rman
- ii) Rman > Configure encryption for db on
- iii) If & backup db plus archivelog.
- iv) Configure encryption Algorithm '1111' for db on;
- v) Set encryption on identified by 'password';
to decrypt backup you need password
- vi) Set decryption identified by 'password';
- vii) Restore db.

Q. How to do restoration? High level of steps restoration
Restoring a db means back to a previous state from a backup.

Steps:

full, incremental, specific backup time stamp.

- ① determine backup :- identify backup you want to restore
- ② prepare for restoration:- ensure that you have necessary files, datafile, control file, archived redo log.
- ③ shut down db.
- ④ copy backup files. → if files not on destination then copy
- ⑤ start db in mount mode.
- ⑥ restore control file
restore control file from 'backlocation';
- ⑦ restore db
- ⑧ recover db
- ⑨ open the db

Q. Suppose restoration is going on & while restoring we need to check redo log file. now, if there are no logfiles what is your action to check logfiles?

- ① check the alert logfile.
- ② check query the v\$log view.
- ③ check the query the v\$logfile view.

If you are using Rman to recovery. recover the file

- ④ restore redo log file

restore logfile 'path' from 'backup_location';

encrypting Rman backups:

- i) Connect with rman
- ii) Rman > Configure encryption for db on
- iii) ~~11~~ & backup db plus archivelog
- iv) configure encryption Algorithm '1111' for db on;
- v) set encryption on identified by 'password';
to decrypt backup you need Password
- vi) Set decryption identified by 'password';
- vii) Restore db.

Q. How to do restoration? High level of steps restoration
Restoring a db means back to a previous state from a backup.

Steps:

- ① determine backup :- identify backup you want to restore
- ② prepare for restoration; ensure that you have necessary files, datafile, control file, archived redo log
- ③ shut down db.
- ④ copy backup files. → if files not on destination then copy
- ⑤ start db in mount mode.
- ⑥ restore control file
- ⑦ restore db
- ⑧ recover db
- ⑨ open the db

Q. Suppose restoration is going on & while restoring we need to check redo log file. now, if there are no logfiles what is your action to check log files?

- ① check the alert log file.
 - ② check query the v\$log view.
 - ③ check the query the v\$logfile view.
- If you are using Rman to recovery - recover the file
- ④ restore redo log file
 - restore Logfile 'path' from 'backup_location';

Q. Explain Startup sequence of a cluster?

Startup

- ① Stop db
- ② Stop ggy
- ③ crsctl stop crs
- ④ If oracle listener is up then stop manually.
- ⑤ Check status.

- Shut down manually
- ① Cluster are up & ready
② If not up then use crsctl start crs
③ Start all listeners
④ Start ggy
⑤ Start db.

Q. Do you know export & import? ~~utility~~

yes.

(index, table, procedure)

Export: is used to extract data & db objects from an oracle db & store them in a binary.

Import: is used to read data & db object from an export dump file created by export utility & load them into an oracle db.

Q. If you are refreshing schema at a time & source name & target name tablespace schema are different, what you do?

① Create mapping betw source & target tablespace.

② Export data with tablespace mapping.

expdp dumpfile=d1.dmp remap_tablespace=source_t1:target_t1, source_t2:target_t2.

③ Transfer & import data:

impdp dumpfile=d1.dmp remap_tablespace=source_t1:target_t1, source_t2:target_t2;

Q. Why remap need to be used?

It allows you to modify object names, tablespace & other attributes during export & import.

Use of remap:- different tablespace during import

① Remap_tablespace:-- allow you to remap object to a

② Remap_schema:--

Q. Suppose you ran import in db & import is running slow, how you can speed up process.

① parallelism :- data Pump Support parallel processing

impdp dumpfile=d1.dmp parallel=4.

② Buffer size :- increase the buffer size.

impdp dumpfile=d1.dmp data_options=Buffer_size: 1024

③ Disable index & constraints :-

impdp dumpfile=d1.dmp transform=disable_constraints, -f indexes,

④ use direct path :-

impdp dumpfile=d1.dmp Direct=Y ... in parameter file

Q. Can I increase filesize while running parallel.

Yes.

When you running data pump operation in parallel you may encounter this scenario.

To increase the file size of each parallel process, you can adjust the filesize parameter.

expdp dumpfile=d1.dmp parallel=4 filesize=500m.

Q. When facing a critical activity at the end of my shift, what will your approach, handover or continue.

→ depend on nature of activity, its urgency & the potential impact on the project. Then I working continue.

Q. What do you think, db activity is small?

No.

db activity is very responsible work, because we handle or check ① transaction volume ② data size

③ complexity of queries. ④ system performance

⑤ Data changes ⑥ application workload

⑦ Business impact.

Q. Tell me any occasion where you taken ownership & completed activity on your own?

Q. what is db link ? How db links works ?
why we used ?
db link is like a virtual bridge that connects different db allowing them to communicate & share create public database link 'dblink_name' connect info. to user identified by 'password' using 'tnsnames.ora' used :

- ① Data integration :- you can fetch & combine data from multiple sources.
- ② cross-db queries :- they enables query that involve tables from different db without the need to physically copy the data.
- ③ data distribution :- db links facilitate distributing data across db.

Q. do you know how to perform dm drill ?
yes,

dm drill is like conducting a practice session to ensure that in the real disaster you can smoothly switch to your backup.

In dm drill activity we perform switch over, switch back & failover. with the help of backup/recreate or recover failover.

for this activity we communicate plan with stakeholders

team members.

After complete the activity collect the feedback & update the D&R Plan.

Q. if you are taking full db backup on your env, how to config. we each db so that size should not go beyond 100 GB.

① we take weekly full backup & other days take incremental backup.

② use backup compression :- Backup as compressed Backupset db plus archive log;

③ delete old archive & old backup in

configure archive log deletion pols to backup 1 time to disk.

yes,
I completed activity own myself regarding
drill like switch over & switch back
on slowness & resolve the issue.

Page No. 08 Date: Youva

Q. Rate yourself in performance tuning out of 10

- 7
- Q. There is a long running query, what is your approach.
 - ① Identify long running query from V\$SESSION, V\$SQL
 - ② Review explain plan: Plan provide insights into how the db is processing the query.
 - ③ Check index usage.
 - ④ Identify locks & blocks are there.
 - ⑤ At the last if issue is not persist then ask to application team can we kill the query if they say yes then we kill query otherwise not.

Q. How to compare that SQL has taken a good plan or bad plan.

Assessing the efficiency of how the db processes the query to retrieve data.

If plan is good

① Low cost (CPU, I/O)

② Index usage

③ Filtering early

④ Avoid full table scan

⑤ Join method optimal

If plan is Bad

① High cost

② Missing index

③ Filtering late

④ Full table scan

⑤ Suboptimal join method

Q. How to check plan hash / plan ID in SQL?

For specific SQL statement by querying the V\$SQL

Eg: Select SQL_ID, plan_hash_value from V\$SQL where

SQL_Text like 'your SQL text';

Q. What is Data Guard?

It is feature provided by Oracle to perform create, manage, disaster recovery, availability, reliability. Standby db that can be used backups in case primary db become unavail.

Key concept of data guard :-

- ① Primary :- it is main prod db. where all the transaction & data modification take place.
- ② Standby :- it is replica of primary db i.e. kept in sync. It is a backup of primary.
- ③ redo logs :- It is the records of changes made to the db. redo logs from the primary db are transmitted to the Standby db to keep it up-to-date.

Modes of operation :

- ① Maximum availability :- both db are synchronized in time. provide highest level of data protection. there might be slight impact on performance due to sync.
- ② Max performance :- it allows for asyn transmission of redo logs to the standby, reduce the impact on the primary db. not take care of standby db. it's fast.
- ③ Max protection :- it ensure that highest level of data protection by requiring acknowledgement from standby db before committing changes on Primary.

Use :-

- ① disaster recovery
- ② planned maintenance
- ③ testing & reporting.

Q. In DR ? How you are protecting data in case you don't have DR ?

Taken regularly backup of our data & ensure that backup copies are stored in secure & separate location from the primary data.

copy of backup store on cloud for security from disaster, earthquake.

From this backup we recover & make db.

Q. How to check sync status of DR ?

Run or check archive log list → Primary.
Select sequence #, applied from V\$Standby_Log
from Standby

Q. what is data masking / data redaction?
it is used to protect sensitive info in db by replacing
the original data.
The goal is to provide a secure environment for testing
development or sharing data without exposing private
Data masking: modifying or hiding sensitive data while
preserving the structure & format of the info.

How data masking works:

- ① identification of sensitive data :- determine which fields
in column in the db contain sensitive info
- ② application of masking rules:-
 - ① Randomization:- Replacing actual value with random
 - ② substitution:- replace value with realistic looking data
 - ③ Encryption:- apply encryption algo to secure data
- ③ integration with application:

Integrate data masking into applications of db so that
the masking rules are automatically applied when
accessing.

Q. what is datafile encryption?

is a security measure that involves encoding the contents of a db datafile to protect sensitive info.

works:-

- ① selection of data to encrypt:- db administrators identify sensitive data
- ② encryption process:- the identified data or datafiles are processed through encryption algo.
- ③ use of encryption key:- it is essentially secret codes, are used during the encryption process.
- ④ storage of encrypted data:- stored in datafile. If someone gain access, the data file without key they will only see scrambled data.

Q. why we need database verbose?

- ① to update registry.

Q. Step by step encryption (datafile, all etc)

- ① connect to the db (with all privileges)
- ② query for tables in a specific tablespace

Select table_name, tablespace_name from all_tables
where tablespace_name = 'User';

- ③ query for columns in a specific table:

Select table_n, column_n, data_type from all_tab_columns
where table_name = 'tablename';

- ④ identify sensitive table.

- ⑤ Enable TDE on the identified tablespace.

- 1) conn / as sysdba —/ with DBA privileges.
- 2) set the wallet & master key.

Administer Key management create Auto-login keystore
From 'wallet path' identified by 'password wallet'.

- 3) enable TDE on the identified tablespace.

Alter tablespace swoopnl encrypt;

- ⑥ reorganize the tables in the encrypted tablespace.

Alter table Tarte move tablespace swoopnl;

- ⑦ Verify encryption in the encrypted tablespace
Select table_n, tablespace_n from all_tables where
table_name = 'tarte'.

- ⑧ Backup the wallet.

Administer Key management set Keystore open by 'wallet-pass'
Container = ALL;

Administer key management backup Keystore to 'path backup'

- ⑨ Restart the db to ensure that changes take effect.

Data Redaction:-

"It is a security feature that dynamically info in the result
set of queries."

rule:-

1) column - select column

2) expression - give condn

3) policy - expression=true.

e.g.; column = salary
expression = ? = HR

Policy : redact the salary

if the user doesn't
have the HR role.

Q. If db of 10tb is compressed what will be the expected size of db backup.
→ 5tb.

Q. RAC architecture:- Real Cluster application.

① Cluster:- It is a key component that manages the cluster environment. provides services such as node membership, communication.

② Cluster interconnect:- is a dedicated, high speed network that enables communication between cluster nodes.

③ Shared Storage:- this shared storage ensures data consistency & allows any instance to read & write to the same files.

④ DB instance:- each node runs its own Oracle DB instance.

⑤ Global Cache Services & Global Enqueue Services:- both are cluster services responsible for managing cache fusion, that ensure that data consistency across multiple instances.

⑥ Cache Fusion:- allows instance to share blocks directly without accessing storage.

⑦ VIP, Scan IP.

Background Processes

① LMS - Lock manager service:- manage keys to prevent 2 guest from entering the same room at the same time.

② Global Cache Service (GCS):- GCS makes sure that instance can share & co-ordinate access to the same data blocks, preventing conflict.

③ Global enqueue service (GES):- ensure that instance can co-ordinate access to resource like locks, preventing multiple instances from interfering with each other.

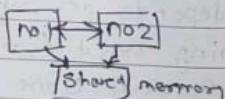
④ Global Cache Resource (GCR):- it keeps track of the data blocks that are being shared among instances in the buffer cache making sure each instance has the most up-to-date used.

→ ① High availability

② Load balancing.

③ Data consistency (cache fusion)

④ Scalability.



- ③ detailed info ④ performance tuning ⑤ Query optimizer
⑥ security auditing.

Date: _____
Year: _____

Q. what is your total experience ?
3 years

Q. why looking for change ?

- ① upgrade the skills ② career growth

Q. what is your current location ?
Pune

Q. what is the size of db in your current project ?
10tb

Q. Do you know Rac environment ?

yes, I'm familiar with Rac.

It is feature of the oracle db that allows a single db to run on multiple servers.

- provide high availability, scalability, load balancing. db provides
- ① Shared storage ② Global cache co-ordination
③ cache fusion ④ ASM

Q. what are technical skills used in day to day activities .

- ① Backup & recovery ② Security management
③ patch management & upgrade ④ db monitoring

Q. How frequent you execute cloning of db ?

It depends on system. like size, performance.

Cloning used for :-

- ① Backup clone ② Test clone ③ development clone
need backup of db with archivelog, pvd, control, sp, p, file.

Q. How db is setup in current project?

There are rac set up of 3 node / 2 node & also have standby standalone db set up also.

Q. How you can start one instance on a host without starting 2nd instance?

sqlplus start instance -d <db_name> -i <instance_name>

Q. sqlplus all command useful:

① sqlplus config database -d <db_name>
displays the configuration info for a rac db, includ. the db name, instance & services.

② sqlplus start database -d <db_name>
start the specified rac db.

③ sqlplus stop database -d <db_name>
stop the specified rac db.

④ sqlplus status database -d <db_name>
displays the status of Specified rac db

⑤ sqlplus config instance -d <db_name> -i <instance_name>

⑥ sqlplus start instance -d <db_name> -i <instance_name>

⑦ sqlplus stop -

⑧ sqlplus status -

⑨ sqlplus config service -d <db_name> -s <service_name>

⑩ sqlplus start service -d -

⑪ sqlplus stop -

⑫ sqlplus status -