

21. How to encrypt the RMAN Backup.

Encrypting RMAN (Recovery Manager) backups in Oracle is crucial for securing sensitive data during the backup process. Oracle provides the ability to encrypt RMAN backups using Transparent Data Encryption (TDE). TDE ensures that the backup files are encrypted, and they remain secure even if the backup files are stored in an unprotected location.

Here's a step-by-step guide on how to encrypt RMAN backups using TDE:

1. Set Up Transparent Data Encryption (TDE):

Before encrypting RMAN backups, you need to set up Transparent Data Encryption for the Oracle database. This involves creating a TDE wallet and enabling encryption for tablespaces or the entire database. Follow the Oracle documentation or your organization's procedures for setting up TDE.

2. Create and Open the TDE Wallet:

Ensure that the TDE wallet is created and opened. The wallet is used to store the master encryption key. You may need to provide a password for the wallet.

sql

Copy code

```
SQL> ALTER SYSTEM SET ENCRYPTION WALLET OPEN IDENTIFIED BY "your_wallet_password";
```

3. Configure RMAN Encryption:

To enable RMAN encryption, you need to set the ENCRYPTION parameter in your RMAN script or directly in the RMAN command line. Specify the algorithm and the encryption key.

sql

Copy code

```
RMAN> SET ENCRYPTION ON IDENTIFIED BY "your_rman_encryption_key" ALGORITHM  
'AES256';
```

The encryption key specified in the RMAN command should match the encryption key used to open the TDE wallet.

4. Perform RMAN Backup:

Run your regular RMAN backup commands with the encryption settings. Here's an example:

sql

Copy code

```
RMAN> BACKUP DATABASE PLUS ARCHIVELOG;
```

The encryption settings specified earlier will apply to this backup operation.

5. Verify the Encrypted Backup:

Once the backup is completed, verify that the backup files are encrypted. You can check the file headers or use Oracle utilities to confirm the encryption status.

Additional Considerations:

Backup Compression: You can combine encryption with RMAN backup compression to

optimize storage space for encrypted backups.

Secure Storage of Encryption Keys: Ensure that the encryption keys (TDE wallet password and RMAN encryption key) are securely stored. Losing access to these keys can result in the inability to restore encrypted backups.

Backup and Recovery Testing: Regularly test the restore process for encrypted backups to ensure that the encryption and decryption mechanisms are working as expected.

By following these steps, you can encrypt RMAN backups using Transparent Data Encryption in Oracle, adding an extra layer of security to your backup strategy. Always refer to Oracle documentation and best practices for Transparent Data Encryption for detailed guidance based on your Oracle database version and deployment.

22.

How to do Restoration? High level steps of restoration?

Certainly! Here's a concise overview of the high-level steps for restoring an Oracle database:

1. Identify Recovery Point:

Determine the specific point in time or SCN (System Change Number) to which you want to restore the database.

2. Ensure Valid Backups:

Verify the availability of valid and up-to-date backups, including full database backups and archived redo log files.

3. Shut Down Database:

Gracefully shut down the Oracle database to prepare for the restoration process.

4. Restore Database Files:

Restore the required database files, including data files, control files, and SPFILE (if applicable), from the backup.

5. Restore Archived Redo Logs:

If necessary, restore archived redo log files generated after the recovery point to enable point-in-time recovery.

6. Create Restore Point (Optional):

Optionally, create a restore point to mark a specific point in the recovery process.

7. Start Database in Mount Mode:

Start the database in mount mode to enable control file restoration and recovery.

8. Restore Control Files:

Restore the control files from the backup to ensure the integrity of the database.

9. Perform Database Recovery:

Apply recovery operations, such as complete recovery or point-in-time recovery, using archived redo logs.

10. Open the Database:

- After successful recovery, open the database for user access.

11. Test Restored Database:

- Conduct tests to validate the restored database's integrity and consistency.

12. Monitor and Resolve Issues:

- Monitor the database for any potential issues post-restoration and address them promptly.

It's essential to note that the steps may vary based on specific Oracle database versions, configurations, and recovery scenarios. Always refer to Oracle documentation and adhere to best practices for database restoration in your environment. Regular testing of backup and recovery procedures is recommended to ensure preparedness for potential disasters.

23. Suppose restoration is going on and while restoring we need to check readlog file. Now, if there are no logfiles, what is your action to check logfiles.

#####

If you are performing a database restoration and need to check the log files (redo log files), but there are no available log files, it's important to take appropriate actions to address the situation. Here are steps you can consider:

Check the Archive Log Destination:

Verify the destination where archived redo log files are supposed to be stored. This is typically defined by the LOG_ARCHIVE_DEST_n initialization parameters.

sql

Copy code

```
SQL> SHOW PARAMETER LOG_ARCHIVE_DEST;
```

Ensure that the archive log destination is accessible and configured correctly.

Check Archive Mode:

Confirm that the database is running in ARCHIVELOG mode. This mode is necessary for the generation and archiving of redo log files.

sql

Copy code

```
SQL> SELECT LOG_MODE FROM V$DATABASE;
```

If the database is not in ARCHIVELOG mode, you may need to switch to this mode to enable the creation of archived redo log files.

Check Archive Process:

Confirm that the ARCH background process is running. This process is responsible for archiving filled redo log files.

sql

Copy code

```
SQL> SELECT PROCESS, STATUS FROM V$MANAGED_STANDBY;
```

Ensure that the ARCH process is active and archiving logs.

Check Log Switches:

Check recent log switches to confirm if the database is actively switching between redo log files.

sql

Copy code

```
SQL> SELECT * FROM V$LOG_HISTORY;
```

If there are no recent log switches, it may indicate that the database is not actively generating redo log files.

Manually Force Log Switch:

You can manually force a log switch to generate a new redo log file. This may be useful during the restoration process.

sql

Copy code

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

After forcing a log switch, check if archived redo log files are generated.

Check File System Space:

Ensure that there is sufficient space on the file system where archived redo log files are stored. Lack of space can prevent the archiving process.

Review Alert Log:

Examine the database alert log for any error messages or warnings related to the archiving process.

sql

Copy code

```
SQL> SHOW PARAMETER BACKGROUND_DUMP_DEST;
```

Locate and review the alert log file in the directory specified by BACKGROUND_DUMP_DEST.

Investigate Error Messages:

If there are error messages related to archiving in the alert log, investigate and address the underlying issues.

By taking these steps, you can identify potential issues preventing the generation or archiving of redo log files during the restoration process. Addressing any identified issues will help ensure the successful completion of the restoration.

24. Explain startup sequence of a cluster.

#####

bash

Copy code

```
srvctl start database -db <DB_NAME>
```

5. Verify Database Status:

Confirm that the Oracle RAC database instances are up and running on all nodes.

bash

Copy code

```
srvctl status database -db <DB_NAME>
```

6. Start Cluster Managed Services (Optional):

If there are additional cluster-managed services configured for the Oracle RAC database (e.g., listeners, VIPs), start them.

bash

Copy code

```
srvctl start service -db <DB_NAME>
```

7. Verify Service Status (Optional):

Confirm that the cluster-managed services are running.

bash

Copy code

```
srvctl status service -db <DB_NAME>
```

8. Start Application Services (Optional):

If there are application services configured and managed by the Oracle Clusterware, start them.

bash

Copy code

```
srvctl start service -db <DB_NAME> -service <SERVICE_NAME>
```

9. Verify Application Service Status (Optional):

Confirm that the application services are running.

bash

Copy code

```
srvctl status service -db <DB_NAME> -service <SERVICE_NAME>
```

10. Perform Functional Testing:

After the startup sequence is completed, perform functional testing to ensure that the Oracle RAC database is accessible, services are responding, and the cluster is operating as expected.

Note: The actual commands and parameters may vary based on the Oracle RAC version, configuration, and naming conventions used in your environment. Additionally, the steps may include additional actions based on the specific features and components employed in your Oracle RAC deployment. Always refer to Oracle documentation and best practices for your specific Oracle RAC version.

25. Do you know export / import.

#####

Yes, I'm familiar with the export and import utilities in Oracle. These utilities are used for logical backups and restoration of Oracle database objects. Here's a brief overview:

1. Oracle Data Pump (expdp / impdp):

Oracle Data Pump, introduced in Oracle Database 10g, is an advanced utility for exporting and importing data and metadata. It provides significant enhancements over the older exp/imp utilities.

Export (Data Pump):

bash

Copy code

```
expdp username/password@connect_string DIRECTORY=dpump_dir1
DUMPFILE=export_file.dmp LOGFILE=export_log.log SCHEMAS=schema_name
```

Import (Data Pump):

bash

Copy code

```
impdp username/password@connect_string DIRECTORY=dpump_dir1
DUMPFILE=export_file.dmp LOGFILE=import_log.log SCHEMAS=schema_name
```

2. Original Export / Import (exp / imp):

The original export and import utilities are still available for backward compatibility, but Oracle recommends using Data Pump for better performance and functionality.

Export (Original):

bash

Copy code

```
exp username/password@connect_string file=export_file.dmp log=export_log.log
```

owner=schema_name
Import (Original):

bash

Copy code

```
imp username/password@connect_string file=export_file.dmp log=import_log.log  
fromuser=schema_name touser=new_schema_name
```

Key Points:

Data Pump Advantages:

Data Pump supports parallelism for faster exports and imports.
It provides more flexible and granular options for data and metadata selection.
Data Pump is the preferred choice for most modern Oracle database versions.
Original Export / Import:

While Data Pump is recommended for new implementations, the original exp/imp utilities may still be used in specific scenarios, especially when migrating from older Oracle versions.

Connect String:

The connect_string typically includes the database connection information, such as hostname, port, and SID (Service Identifier).

Directory:

The DIRECTORY parameter specifies the directory object in the database where the export or import files are located.

Dumpfile and Logfile:

The DUMPFILE parameter specifies the name of the export file, and LOGFILE specifies the name of the log file.

Remember to replace placeholders like username, password, connect_string, schema_name, new_schema_name, export_file.dmp, export_log.log, and dpump_dir1 with your actual values.

Always refer to the Oracle documentation for the specific version you are using for detailed and up-to-date information on export and import utilities.

26. If you are refreshing 5 schema at a time and source name & target name tablespace schema are different.

#####

When refreshing multiple schemas, each with different tablespace names, from a source environment to a target environment, there are several considerations to

ensure a successful and accurate schema refresh. Here's a high-level guide:

1. Export Source Schemas:

Use Oracle Data Pump (expdp) or the original export (exp) utility to export the data and metadata of the source schemas. Specify the tablespaces and other relevant options during the export.

bash

Copy code

```
expdp username/password@source_db schemas=schema1,schema2,schema3
directory=dpump_dir1 dumpfile=export_file.dmp logfile=export_log.log
```

2. Import into Target Schemas:

Import the exported dump file into the target schemas in the target database. Ensure that you map the tablespaces appropriately during the import to match the target environment.

bash

Copy code

```
impdp username/password@target_db remap_schema=schema1:schema1_target
remap_tablespace=source_ts1:target_ts1 directory=dpump_dir1
dumpfile=export_file.dmp logfile=import_log.log
Repeat this step for each schema you are refreshing.
```

3. Adjust Constraints and Indexes:

If there are referential constraints or indexes that reference tables across different tablespaces, you may need to adjust them after the import to reflect the new tablespaces in the target environment.

4. Verify Data and Objects:

After the import, thoroughly verify the data, constraints, indexes, and other database objects in the target schemas to ensure they match the source schemas.

5. Update Statistics:

Gather or update statistics for the refreshed schemas to ensure the query optimizer has accurate information for optimization.

6. Resolve Dependencies:

If there are dependencies between the refreshed schemas, ensure that these dependencies are resolved. This may include stored procedures, triggers, or any other database objects that reference objects in other schemas.

7. Review and Adjust Privileges:

Review and adjust the privileges and roles for the users in the target schemas based on the requirements of the application.

8. Testing:

Conduct thorough testing of the refreshed schemas to validate the integrity and functionality of the data and applications.

Note:

The `remap_schema` and `remap_tablespace` parameters in the Data Pump import command are crucial for remapping schemas and tablespaces during the import process. Adjust the commands based on your specific requirements, tablespaces, and schema names.

Always perform schema refresh activities in a test or staging environment before applying them to a production environment.

Remember to adapt the provided commands based on your actual schema and tablespace names in both the source and target environments. Additionally, consult Oracle documentation and best practices for more detailed guidance on schema refresh operations.

27. Why is remap need to be used?

The REMAP functionality in Oracle Data Pump is used to alter or remap certain attributes during the import process. The primary use case for the REMAP feature is to address differences between the source and target databases, ensuring a smooth import when specific conditions need to be adjusted. Here are some common scenarios where REMAP is beneficial:

Schema Name Change:

When importing data into a target database, you may want to use a different schema name for the imported data than what was used in the source. The `REMAP_SCHEMA` option allows you to remap the schema name during the import.

bash

Copy code

```
impdp username/password@target_db REMAP_SCHEMA=source_schema:target_schema ...
```

Tablespace Name Change:

If the tablespaces in the target database have different names than those in the source, you can use `REMAP_TABLESPACE` to remap the tablespaces during the import.

bash

Copy code

```
impdp username/password@target_db REMAP_TABLESPACE=source_ts:target_ts ...
```

Tablespace Objects Relocation:

When you want to move objects (tables, indexes, etc.) to different tablespaces in the target database, `REMAP_TABLESPACE` can be used to remap specific objects.

bash

Copy code

```
impdp username/password@target_db  
REMAP_TABLESPACE=source_ts1:target_ts1,source_ts2:target_ts2 ...
```

Changing Storage Characteristics:

You may want to change storage characteristics (such as `INITIAL` and `NEXT` extents)

during the import. The REMAP_TABLE option can be used for this purpose.

bash

Copy code

```
impdp username/password@target_db REMAP_TABLE=source_table:target_table ...
```

Changing Directory Paths:

If the directory paths used in the source and target databases differ, the REMAP_DATAFILE option can be used to remap the directory paths for datafiles.

bash

Copy code

```
impdp username/password@target_db REMAP_DATAFILE=source_path:target_path ...
```

Global Object Name Changes:

When importing into a different database with a different global database name, REMAP_GLOBAL_NAME can be used to remap global object names.

bash

Copy code

```
impdp username/password@target_db  
REMAP_GLOBAL_NAME=source_global_name:target_global_name ...
```

In summary, the REMAP options in Oracle Data Pump provide flexibility during the import process, allowing you to adapt the import to the specific characteristics and requirements of the target database. It's a powerful feature for handling differences between source and target environments, ensuring a successful and accurate data import. Always refer to the Oracle documentation for the specific version you are using for detailed information on REMAP options and their usage.

28. Suppose you ran import in db and import is running slow, how you can speedup process.

If an Oracle Data Pump import is running slow, there are several strategies you can employ to improve the performance of the import process. Here are some tips:

Parallelism:

Use the PARALLEL parameter to increase the degree of parallelism during import. This can significantly speed up the import process by utilizing multiple worker processes.

bash

Copy code

```
impdp username/password@target_db PARALLEL=4 ...
```

Increase Buffer Size:

The DATA_BUFFER parameter controls the size of the buffer used for reading data during import. Increasing the buffer size can enhance performance.

```
bash
Copy code
impdp username/password@target_db DATA_BUFFER=2M ...
Commit Interval:
```

The COMMIT_INTERVAL parameter specifies the number of rows after which a commit is performed. Increasing the commit interval can reduce the overhead of frequent commits.

```
bash
Copy code
impdp username/password@target_db COMMIT_INTERVAL=10000 ...
Direct Path Load:
```

Use the DIRECT parameter to enable direct path loading, which can be faster than conventional path loading. However, direct path loading may have certain restrictions.

```
bash
Copy code
impdp username/password@target_db DIRECT=Y ...
Disable Constraints and Indexes:
```

Temporarily disable constraints and indexes during the import and enable them afterward. This can accelerate the import process.

```
bash
Copy code
impdp username/password@target_db TRANSFORM=DISABLE_ARCHIVE_LOGGING:Y ...
Use Network Link for Remote Imports:
```

If importing over a network link, consider using the NETWORK_LINK parameter for direct import between databases.

```
bash
Copy code
impdp username/password@target_db NETWORK_LINK=source_db ...
Monitoring and Tuning:
```

Monitor the import progress using the Data Pump status and log files. If specific tables or objects are causing performance issues, you might consider excluding or adjusting them.

Optimize Storage:

Ensure that the tablespaces and datafiles in the target database are appropriately sized and configured to accommodate the imported data.

Database Parameters:

Adjust Oracle database parameters such as DB_BLOCK_SIZE, SORT_AREA_SIZE, and PGA_AGGREGATE_TARGET to optimize memory usage during import.

Indexes:

Consider importing indexes after the data to speed up the initial data load. You can use the EXCLUDE=INDEX option during the initial import.

bash

Copy code

```
impdp username/password@target_db EXCLUDE=INDEX ...
```

Always consider the specific characteristics of your environment and the nature of the data being imported when applying these tips. Experimentation and testing in a non-production environment are recommended to find the optimal settings for your particular use case.

29. can I increase filesize in paramater file while running parallel.

Copy code

```
impdp username/password@target_db PARALLEL=4 DIRECTORY=dpump_dir1
```

```
DUMPFIL=exp%U.dmp FILESIZE=1G ...
```

In this example:

PARALLEL=4 indicates that the import will run with four parallel processes.

DIRECTORY=dpump_dir1 specifies the directory object in the database where the dump files will be written.

DUMPFIL=exp%U.dmp is a pattern for generating unique dump file names for each parallel process.

FILESIZE=1G sets the maximum size for each dump file to 1 gigabyte.

Adjust the value of FILESIZE based on your requirements and the capacity of your storage system. It's important to note that the total disk space required for the import will be the product of the FILESIZE and the number of parallel processes.

Always consider the available disk space and the performance characteristics of your storage subsystem when determining the optimal value for FILESIZE.

Additionally, monitor the import progress and adjust parameters as needed to optimize performance.