

32. Tell me any occasion where you taken ownership and completed activity on your own.

#### ##### Database Performance Optimization

There is an Oracle database that has been experiencing performance issues. Users have reported slow response times, and it's affecting critical business operations. In this situation, I am (DBA) takes ownership of the performance optimization initiative.

##### Identification of Performance Issues:

The DBA conducts a thorough analysis of the database using performance monitoring tools, Oracle Enterprise Manager, or other diagnostic queries. They identify specific SQL queries, inefficient indexes, or other bottlenecks contributing to the performance degradation.

##### Development of Optimization Plan:

Based on the analysis, the DBA formulates a comprehensive plan to optimize the database. This plan may include rewriting SQL queries for better efficiency, creating or modifying indexes, adjusting configuration parameters, or redistributing data across tablespaces.

##### Coordination with Stakeholders:

The DBA communicates with relevant stakeholders, such as application developers, system administrators, and business users, to discuss the optimization plan, potential downtime, and the expected impact on system performance during the optimization process.

##### Implementation of Changes:

Taking ownership of the optimization process, the DBA implements the planned changes during a scheduled maintenance window. This may involve executing SQL scripts, modifying database structures, or adjusting configuration settings.

##### Monitoring and Testing:

After the changes are implemented, the DBA closely monitors the database performance in real-time. They conduct testing to ensure that the optimizations have positively impacted response times and have not introduced any adverse effects.

##### Documentation and Reporting:

The DBA documents the optimization activities performed, including changes made, reasons for those changes, and the observed improvements in performance. This

documentation serves as a reference for future troubleshooting and analysis.

#### Continuous Improvement:

The DBA remains vigilant about the database's performance, continuously monitoring and addressing any emerging issues. They may also proactively suggest long-term strategies for maintaining optimal database performance.

In this scenario, the DBA takes ownership of the performance optimization initiative, demonstrating a proactive and responsible approach to ensuring the smooth operation of the Oracle database.

33. Rate yourself in performance tuning out of 10.

##### 7/10

34. There is a long running query, what is your approach to check ?

##### When dealing with a long-running query, it's important to identify the cause of the delay and optimize the query for better performance. Here's a systematic approach to check and address a long-running query:

Identify the Long-Running Query for v\$sqlsession\_longops view

```
set linesize 400 pagesize 300
col sid format a10
col SERIAL# format a10
col opname format a20
col target format a30
col sofar format a20
col totalwork format a10
select sid, SERIAL#, username, opname, target, sofar, totalwork,
(totalwork-sofar)/time_remaining bps, time_remaining/60 rmt,
sofar/totalwork*100 fertig, sql_id from v$sqlsession_longops where time_remaining > 0
order by 8;
```

Examine query and Execution Plan

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY_CURSOR('<SQL_ID>'));
```

Check Indexes Ensure that the tables involved in the query have appropriate indexes. Missing or inadequate indexes can significantly impact query performance. Use tools like Oracle's SQL Tuning Advisor to suggest index improvements.

Review Query Structure:

Evaluate the structure of the query. Ensure that it is written efficiently, uses appropriate joins, and includes only the necessary columns and conditions. Avoid using wildcard selects (SELECT \*) and consider limiting the result set.

check table stale and fragmentation

```
SELECT table_name, last_analyzed, stale_stats
FROM dba_tab_statistics
WHERE table_name = 'employee';
```

then statgather or reorg in this table

35. How to compare that sql has taken a good plan or bad plan?

##### Evaluating whether a SQL query has a good or bad execution plan involves assessing the efficiency of the plan in terms of resource utilization, response time, and overall query performance.

Execution Time:

Measure the actual time taken by the query to execute. A good execution plan should result in a reasonably fast execution time. You can compare the execution time against historical values or predefined performance targets.

Resource Utilization:

Evaluate the resource consumption during query execution, including CPU usage, memory, and I/O operations. A good execution plan minimizes resource utilization, avoiding unnecessary strain on the database server.

#### Query Cost:

Many database management systems provide a query cost or estimated cost associated with an execution plan. Lower query costs generally indicate more efficient plans. Review the cost values provided in the execution plan details.

#### Access Methods:

Assess the access methods employed by the execution plan. Look for efficient index usage, full table scans when appropriate, and appropriate join methods. The plan should leverage indexes and minimize data scans where possible.

#### Filtering and Predicate Pushdown:

Check if filters and predicates are applied early in the execution process. A good plan pushes down filtering operations to reduce the amount of data processed during execution.

#### Parallel Execution:

In some cases, parallel execution of a query can enhance performance. Evaluate whether the execution plan utilizes parallel processing efficiently, especially for resource-intensive queries.

#### Index Usage:

Confirm that indexes are used appropriately. A good execution plan utilizes indexes for selective queries and avoids unnecessary full table scans. Ensure that the index statistics are up-to-date.

#### Join Methods:

Examine the join methods employed in the plan. Nested loop joins, hash joins, and merge joins have different performance characteristics. The chosen join method should be suitable for the data distribution and join conditions.

#### Table and Index Statistics:

Ensure that the statistics for tables and indexes are accurate and up-to-date. Outdated statistics can lead to suboptimal execution plans. Regularly gather or update statistics using tools like Oracle's DBMS\_STATS package.

36.How to check plan hash / plan ID in sql?

##### To check the plan hash or plan ID for a SQL query in Oracle, you can query the V\$SQL or V\$SQL\_PLAN views. These views provide information about executed SQL statements and their execution plans. The plan hash value uniquely identifies a specific execution plan for a SQL statement. Here's how you can check the plan hash or plan ID:

```
SELECT sql_id, plan_hash_value, sql_text
FROM v$sql
WHERE sql_text LIKE 'Your SQL Query%';
```

or

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY_CURSOR('<SQL_ID>'));
```

38. What is dataguard?

##### Data Guard is a feature in Oracle Database that provides high availability, data protection, and disaster recovery for enterprise databases. It is a solution for creating and maintaining standby databases, which are synchronized copies of the primary database. The primary database and its standby(s) operate in a coordinated fashion to ensure data redundancy and continuity of operations. Here are key aspects of Oracle Data Guard:

Primary Database:

The primary database is the main operational database that handles read and write operations. It serves as the primary source of data.

Standby Database:

The standby database is a synchronized copy of the primary database. It can be located at a remote site for disaster recovery purposes. There are different types of standby databases, including physical standby, logical standby, and snapshot standby, each with specific use cases.

Redo Transport:

Redo data, generated by transactions on the primary database, is continuously transmitted to the standby database(s). This ensures that the standby databases are kept up-to-date with changes made on the primary database.

Apply Mechanism:

Standby databases apply the redo data received from the primary database to keep their data in sync. The apply mechanism depends on the type of standby database:

Physical Standby: Redo data is applied at the block level.

Logical Standby: SQL statements are transformed and applied to maintain logical consistency.

Snapshot Standby: Acts as a read/write standby, allowing temporary divergence from the primary database.

Switchover:

Switchover is a planned, controlled role transition between the primary and standby databases. It is often used for maintenance tasks, upgrades, or to temporarily make a standby database the new primary.

Failover:

In the event of a failure or outage affecting the primary database, failover allows a standby database to be quickly and automatically promoted to become the new primary database. This helps minimize downtime.

Data Protection and Disaster Recovery:

Data Guard provides data protection by maintaining a synchronized copy of the data at a remote location. In case of a disaster or primary database failure, the standby database(s) can be activated to take over operations.

Role Transitions:

Role transitions involve changing the role of a database from primary to standby or vice versa. These transitions can be planned (switchover) or unplanned (failover).

Data Guard Broker:

Data Guard Broker is a management and monitoring tool that simplifies the configuration, monitoring, and maintenance of Data Guard environments. It provides a centralized interface for managing multiple databases in a Data Guard configuration.

Maximum Availability Architecture (MAA):

Data Guard is a key component of Oracle's Maximum Availability Architecture, a set of best practices and technologies designed to achieve the highest levels of availability and reliability for Oracle databases.

Data Guard is widely used in mission-critical environments where uninterrupted database operations, data protection, and disaster recovery are crucial requirements.

39.In DR? - How you are protecting data in case you dont have DR?

##### Regular Backups:

Implement a robust backup strategy. Regularly back up your data and ensure that backup copies are stored in a secure and separate location from the primary data. This can help recover data in case of accidental deletions, data corruption, or localized failures.

#### Offsite Storage:

Store backup copies in an offsite location to protect against site-wide disasters, such as fires, floods, or earthquakes. Having backups in a geographically different location reduces the risk of losing data due to a localized event.

#### Cloud Storage:

Consider leveraging cloud storage services for backups. Cloud providers often have redundant data centers across different regions, offering a level of geographical separation and ensuring data availability even in the face of regional disasters.

#### Versioning and Snapshots:

Use versioning or snapshots to maintain multiple versions of your data. This can be helpful in recovering from accidental changes or data corruption by reverting to a known good state.

#### Database Replication:

If possible, implement database-level replication features provided by your database management system. This can create copies of the database on separate servers, offering some level of redundancy.

#### High Availability (HA) Solutions:

Explore high availability solutions that provide automated failover in case of hardware or software failures. While not a complete DR solution, HA solutions can minimize downtime and data loss.

#### Redundant Systems:

Deploy redundant systems and components for critical infrastructure. Redundancy in servers, storage, and network components can reduce the impact of hardware failures.

#### Business Continuity Planning:

Develop and document a business continuity plan that outlines procedures for data recovery and continuity of operations in the event of disruptions. This includes communication plans, roles and responsibilities, and specific steps to be taken during incidents.

#### Security Measures:

Implement strong security measures to protect against data breaches and unauthorized access. Regularly update and patch systems to address security vulnerabilities.

#### Monitoring and Alerts:

Implement monitoring systems that can detect anomalies or issues with data integrity. Set up alerts to notify administrators of potential problems so that

they can be addressed promptly.

While a comprehensive Disaster Recovery plan is ideal for ensuring data protection in various scenarios, implementing the above practices can enhance your organization's ability to recover data and maintain business continuity even in the absence of a dedicated DR setup. It's important to tailor these measures to the specific needs and risks of your organization.

#### 40. How to check sync status of DR?

##### To check the synchronization status of a Data Guard configuration in Oracle, you can use the Data Guard broker or query specific views in the database. Here are a few methods to check the synchronization status:

Using Data Guard Broker:

Connect to the Primary Database:

Connect to the primary database using SQL\*Plus or another SQL client.

Start Data Guard Broker:

Start the Data Guard broker by executing the following SQL command:

sql

Copy code

```
ALTER SYSTEM SET DG_BROKER_START=TRUE;
```

Query Data Guard Status:

Use the following command to query the Data Guard configuration status:

sql

Copy code

```
SELECT * FROM V$DATAGUARD_STATUS;
```

This query will provide information about the synchronization status, including the state of the primary and standby databases.

Query Detailed Configuration:

For more detailed information, you can query the V\$DATABASE view on both the primary and standby databases:

sql

Copy code

```
SELECT NAME, OPEN_MODE, DATABASE_ROLE FROM V$DATABASE;
```

This query will show the current role (PRIMARY or STANDBY) and the open mode of the database.

Using Data Guard Views:

Query V\$DATAGUARD\_STATUS:



On the primary or standby database, you can use the following query to check the Data Guard status:

sql

Copy code

```
SELECT * FROM V$DATAGUARD_STATUS;
```

This query provides information about the current state of the Data Guard configuration.

Query V\$ARCHIVE\_DEST\_STATUS:

To check the archiver process status on the primary and standby databases, you can use the following query:

sql

Copy code

```
SELECT DEST_NAME, STATUS FROM V$ARCHIVE_DEST_STATUS;
```

This query shows the status of the archiver process, which is crucial for maintaining synchronization.

Query V\$LOG:

To check the log sequence numbers on the primary and standby databases, use the following query:

sql

Copy code

```
SELECT GROUP#, THREAD#, SEQUENCE# FROM V$LOG;
```

Compare the log sequence numbers between the primary and standby databases to ensure synchronization.

These queries provide real-time information about the Data Guard synchronization status. Ensure that the primary and standby databases are reachable and that the necessary privileges are granted to execute these queries. Monitoring the synchronization status regularly is crucial for maintaining a reliable and effective Data Guard configuration.

41. What is DB Link ? How DB Link works ?

##### A database link (DB Link) in the context of databases, particularly Oracle databases, is a feature that allows you to connect and access objects (tables, views, etc.) in one database from another. It enables communication and data retrieval between databases, even if they are on different servers or instances. Here's how a database link works:

Definition of Database Link:

A database link is a named connection that is defined in the Oracle database. It includes information such as the network address, authentication details, and the

remote database to which it connects.

Creation of Database Link:

To create a database link, a user with the necessary privileges executes a CREATE DATABASE LINK SQL statement. This statement specifies the connection details and credentials for the remote database.

sql

Copy code

```
CREATE DATABASE LINK remote_db
CONNECT TO username IDENTIFIED BY password
USING 'remote_database_tns';
```

In this example, remote\_db is the name of the database link, username and password are the credentials for the remote database, and remote\_database\_tns is the TNS (Transparent Network Substrate) entry for the remote database.

Accessing Remote Objects:

Once the database link is created, users can reference objects in the remote database using a special syntax. For example, to query a table named employees in the remote database, you would use the following syntax:

sql

Copy code

```
SELECT * FROM employees@remote_db;
```

Here, remote\_db is the name of the database link created earlier.

Security Considerations:

When creating a database link, it's important to consider security. The credentials used in the CREATE DATABASE LINK statement determine the level of access the connected user has in the remote database. Users must have the necessary privileges on the remote objects they are trying to access.

Types of Database Links:

There are two main types of database links in Oracle: private and public. Private database links are specific to a user and are only accessible by that user. Public database links are shared among multiple users.

Transaction Control:

Transactions involving a database link are subject to commit and rollback operations. If a transaction is committed on the local database, the changes affecting the remote database through the link are also committed. Similarly, a rollback on the local database also rolls back the changes on the remote database.

Network Communication:

When a query is executed across a database link, the Oracle database establishes a connection to the remote database, sends the SQL statement for execution, and retrieves the results. Network communication plays a crucial role in the performance of queries involving database links.

Performance Considerations:

While database links provide a convenient way to access data across databases, performance considerations, such as network latency and the volume of data being transferred, should be taken into account. Indexing and optimizing queries involving remote objects can help improve performance.

Database links are useful in scenarios where data is distributed across multiple databases, and there is a need to query or update data across these databases seamlessly. They provide a means for building distributed database architectures and facilitating data integration across an organization's information landscape.

42. Do you know how to perform DR Drill?

##### Performing a Disaster Recovery (DR) drill is a crucial activity to test the readiness and effectiveness of your DR plan. A DR drill simulates a real disaster scenario to ensure that the organization can successfully recover its IT systems and data in the event of a disaster. Here's a general guide on how to perform a DR drill:

Pre-Drill Preparation:

Define Objectives:

Clearly define the objectives and scope of the DR drill. Outline the specific systems, applications, and services that will be included in the simulation.

Notify Stakeholders:

Inform all relevant stakeholders, including IT teams, business continuity teams, and external partners, about the upcoming DR drill. Clearly communicate the schedule and expected duration.

Backup Systems:

Take backups of critical systems and data before initiating the DR drill. This ensures that you have a known good state to revert to if needed.

Documentation Review:

Review and update documentation related to the DR plan. Ensure that all procedures, contact lists, and recovery steps are up-to-date.

DR Drill Execution:

Initiate the Simulation:

Start the DR drill by simulating the triggering event, whether it's a data center outage, hardware failure, or other disaster scenario.

Activate DR Plan:

Execute the DR plan as if it were a real disaster. This includes activating standby systems, restoring data from backups, and implementing any necessary recovery procedures.

Monitor and Evaluate:

Monitor the progress of the DR drill in real-time. Evaluate the effectiveness of each step, including the time taken to complete tasks and the overall success of the recovery process.

Communication:

Communicate regularly with stakeholders during the DR drill. Provide status updates, and if applicable, inform users and customers about the temporary unavailability of services.

Inject Realism:

Inject realism into the drill by introducing unexpected challenges or variations. This could include simulated network issues, additional failures, or changes in the scenario to test the adaptability of the DR plan.

Data Validation:

Validate the recovered data to ensure its integrity. Perform checks to verify that the recovered systems and applications are functioning correctly.

Post-Drill Activities:

Debriefing:

Conduct a debriefing session with key participants and stakeholders. Discuss the outcomes, challenges faced, and lessons learned during the DR drill.

Documentation Update:

Update the DR plan documentation based on the findings and feedback from the drill. Make any necessary adjustments to improve the plan's effectiveness.

Training and Awareness:

Provide additional training or awareness sessions for staff based on the lessons learned during the drill. This helps in enhancing the skills and readiness of the DR team.

Report and Analysis:

Generate a comprehensive report that includes a summary of the DR drill, observed performance metrics, and recommendations for improvement. Use this report for analysis and as a basis for refining the DR plan.

Iterative Improvement:

Use the insights gained from the DR drill to iteratively improve the DR plan.

Address any identified weaknesses or areas for improvement to enhance the organization's overall readiness for a real disaster.

Performing regular DR drills is essential for validating the effectiveness of your DR plan, identifying areas for improvement, and ensuring that your organization can respond effectively to unforeseen events.

PRIMARY

```
select switchover_status from v$database;  
select name,open_mode,log_mode,database_role from v$database;
```

```
show parameter log_archive_dest_state;
archive log list;
```

STANDBY

```
select name,open_mode,log_mode,database_role from v$database;
select sequence#,first_time,next_time,applied from v$archived_log order by
sequence#;
```

PRIMARY

```
alter system switch logfile;
alter database commit to switchover to physical standby;
shut immediate
startup mount
select name,open_mode,log_mode,database_role from v$database;
```

STANDBY

```
select name,open_mode,log_mode,database_role from v$database;
alter database recover managed standby database cancel;
select process,sequence#,status from v$managed_standby;
select switchover_status from v$database;
```

PRIMARY

```
alter database recover managed standby database disconnect from session;
select process,sequence#,status from v$managed_standby;
```

STANDBY

```
select switchover_status from v$database;
alter database commit to switchover to primary with session shutdown;
alter database open;
select name,open_mode,log_mode,database_role from v$database;
archive log list;
alter system switch logfile;
/
/
archive log list;
```

PRIMARY

```
select sequence#,first_time,next_time,applied from v$archived_log order by
sequence#;
```

